

# Training : Computer Vision

Hackathon CentraleSupélec-ESSEC

14/02/2022

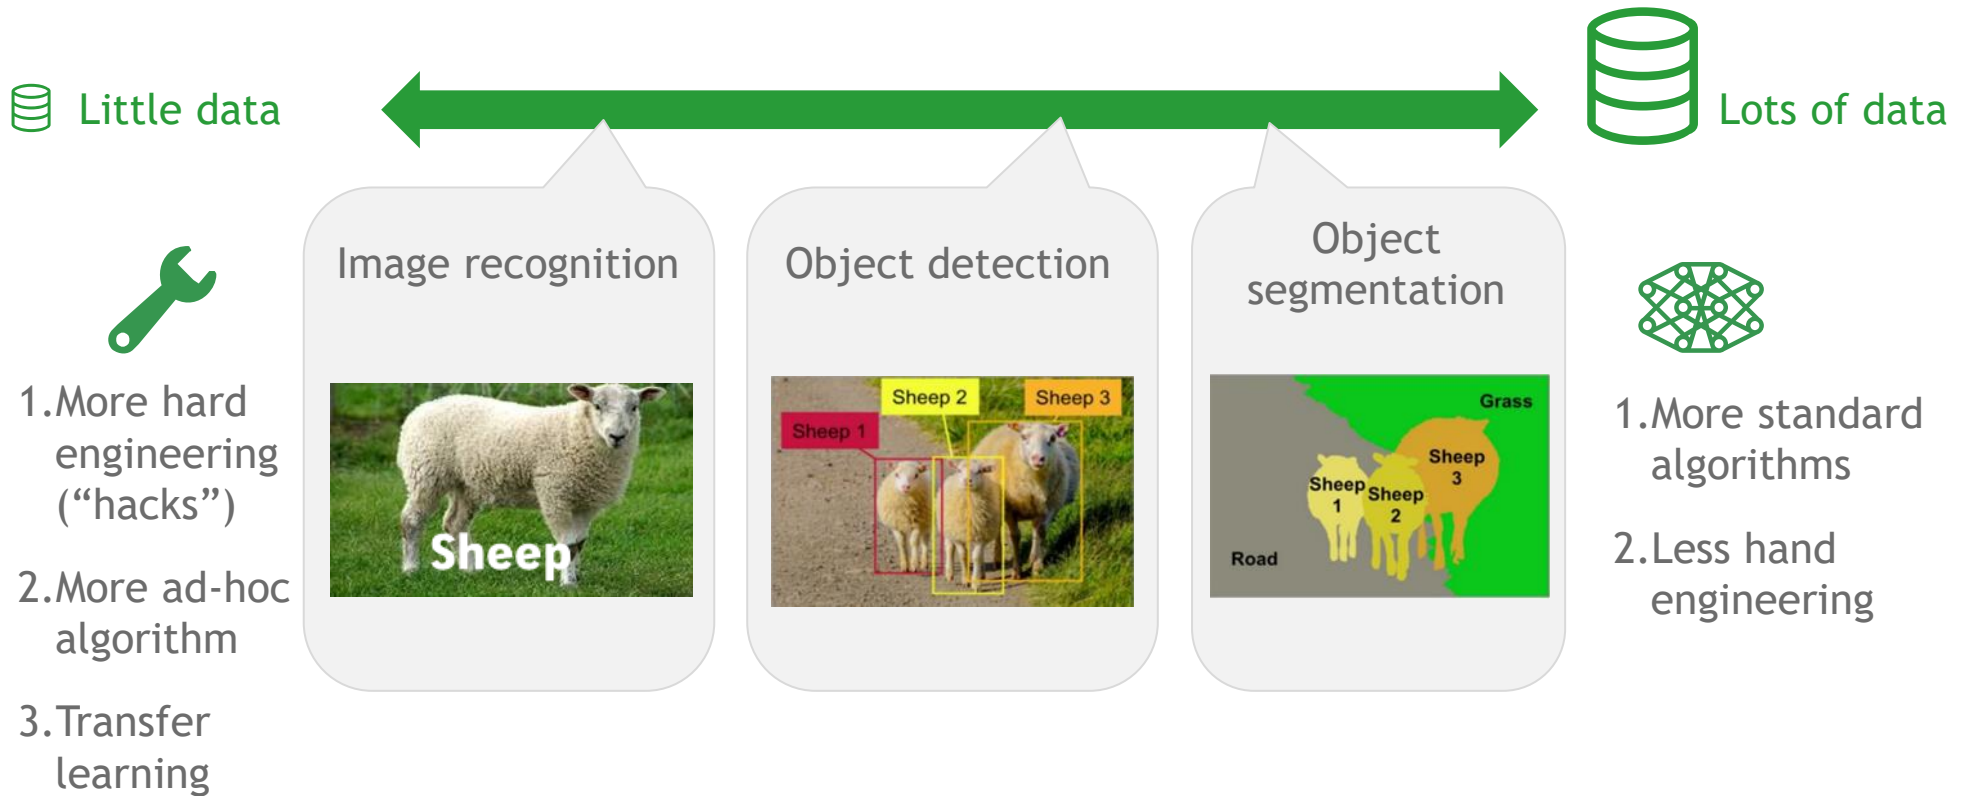


# AGENDA



1. **Computer vision approaches**
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm

The **amount of data available** is a critical element to consider in an image-based application



### Sources of knowledge

1. Labeled data (e.g. image → dog & cat)
2. Hand engineered features, network architecture, etc.

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

## Image Pre-processing


Image scaling  
Contrast  
enhancement  
Cropping and  
padding  
Noise reduction



(a)



(b)



(c)



(d)

(a)Original image

(b)Resized

(c)Grayed

(d)Blured

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)





### Image Pre-processing

Image scaling  
Contrast enhancement  
Cropping and padding  
Noise reduction



### Feature extraction

Lines/Edge detection  
Corners/Blobs/Points  
(SIFT/SURF, PCA, Watershed)





Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

### Image Pre-processing

Image scaling  
Contrast  
enhancement  
Cropping and  
padding  
Noise reduction



### Feature extraction

Lines/Edge detection  
Corners/Blobs/Points  
(SIFT/SURF, PCA,  
Watershed)



### Classification algorithm

CNN



Cat

Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

## Image Pre-processing

Image scaling  
Contrast enhancement  
Cropping and padding  
Noise reduction



## Feature extraction


Lines/Edge detection  
Corners/B  
(SIFT/SU  
Water



## Classification algorithm

## Object detection and/or Identification

YOLO



Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (1/2)

### Image Pre-processing

Image scaling  
Contrast  
enhancement  
Cropping and  
padding  
Noise reduction




### Feature extraction

Lines/Edge detection  
Corners/Blobs/Points  
(SIFT/SURF, PCA,  
Watershed)



### Segmentation

Faster R-CNN





Computer vision can be broken down into a **multitude of blocks** that can be combined as needed (2/2)

### Image Pre-processing

Image scaling  
Contrast enhancement  
Cropping and padding  
Noise reduction



### Feature extraction

Lines/Edge detection  
Corners/Blobs/Points  
(SIFT/SURF, PCA, Watershed)



### Classification algorithm

CNN



### Object detection and/or Identification

YOLO




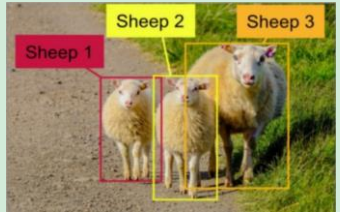
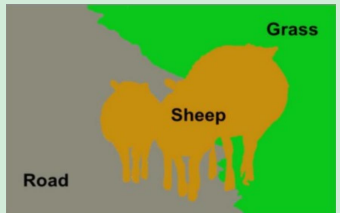

- We can start our process by applying image pre-processing (for instance if we have few data)
- Later we can use detection algorithm in order to find the object of interest
- Finally we can leverage feature extraction to post process the results to make them more precise or follow some field rules

# AGENDA



1. Computer vision approaches
2. **An overview of the main Deep Learning algorithms**
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm

# The 3 main approaches to machine learning vision are Classification, Detection and Segmentation

✂️ Approach	🎯 Objective	📁 Algorithms	👤 Examples
<b>Classification</b>	Associate a class, defining its content, to each image	<ul style="list-style-type: none"> <li>CNN</li> </ul>	
<b>Detection</b>	Spot objects in an image according to predefined categories	<ul style="list-style-type: none"> <li>YOLO</li> <li>SSD</li> <li>Faster R-CNN</li> <li>RFCN</li> </ul>	
<b>Segmentation</b>	<b>Semantic:</b> Link a category to each pixel in an image. Recognize a set of pixels that form distinct categories	<ul style="list-style-type: none"> <li>Unet</li> <li>PSPNet</li> <li>ICNet</li> <li>DeepLab</li> </ul>	
	<b>Per instance:</b> Associate each pixel with the instance to which it belongs.	<ul style="list-style-type: none"> <li>Mask R-CNN</li> <li>ResNet</li> </ul>	

These different approaches are based on **two different types of annotation**




### Polygon

- Polygon annotations select all pixels belonging to an instance. It is used in segmentation



### Bounding box

- Bounding box annotations give position, size and type of object contained in a rectangle parallel to the axes: used in detection






# AGENDA



1. Computer vision approaches
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. Measure the performance to improve your algorithm

To train these models several **datasets** are available online

 Name	 Type	 Annotation
COCO	Natural landscapes	Instance segmentation
ADE20K		
SUN		Bounding boxes detection
IMAGE NET		
PASCAL VOC2012	City landscapes	Instance segmentation
CITYSCAPES		

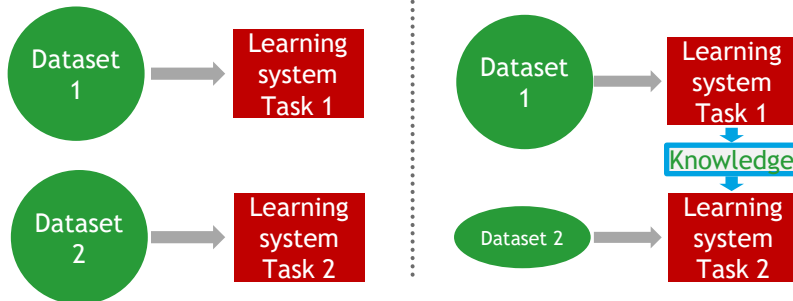
- Is the **context** of the dataset close to context of my images?
- Does the **format** of the dataset compatible with my algorithm?
- What **classes** are labeled in the dataset?
- Is the dataset **balanced** between classes?



# Transfer learning uses knowledge acquired for one task to solve related ones

## General overview


- **Isolated**, single task learning:
  - Knowledge is not retained or accumulated
  - Learning is performed without considering past learned knowledge in other tasks
- **Learning** of a new task relies on the **previous learned tasks**:
  - Learning process can be faster, more accurate and/or need less training data



- There are two **variations** of transfer learning:
  - **Same domain but different tasks**
  - **Same task, but different domains**

## Transfer learning in deep learning

- Use **pre-trained model** and **freeze** most layers of the neural network
  - Freeze layer imply that **weights are not updated** during the backpropagation, the learning rate (LR) is equal to 0
- Apply **fine tuning** to **lasts layers** in order to adapt the model to the specific task
  - The **weights** of fine-tuned layers are **updated** during the backpropagation



- **Keep** simple and **general features** from shallow layers
- **Retrain** complex features from deep layers

# AGENDA



1. Computer vision approaches
2. An overview of the main Deep Learning algorithms
3. The “open” dataset : our gold mine
4. **Measure the performance to improve your algorithm**

Several **metrics** can be used in computer vision depending on the objectives to reach

## ⚙️ Local Metrics


### ▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\mathbf{Box}_{\text{True}} \cap \mathbf{Box}_{\text{Predicted}})}{\text{aire}(\mathbf{Box}_{\text{True}} \cup \mathbf{Box}_{\text{Predicted}})}$$

We set a threshold  $\varepsilon \in [0,1]$

Prediction are considered true if  $IoU \geq \varepsilon$ .



## Model Metrics

# Several **metrics** can be used in computer vision depending on the objectives to reach

## ⚙️ Local Metrics


### ▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$

We set a threshold  $\varepsilon \in [0,1]$

Prediction are considered true if  $IoU \geq \varepsilon$ .



## Model Metrics

### ▪ Mean Intersect over Union (mIoU)

For a fixed **minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: on/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret

# Several **metrics** can be used in computer vision depending on the objectives to reach

## ⚙️ Local Metrics


### ▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$

We set a threshold  $\varepsilon \in [0,1]$

Prediction are considered true if  $IoU \geq \varepsilon$ .



### ▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**

$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**



## Model Metrics

### ▪ Mean Intersect over Union (mIoU)

For a fixed **minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: on/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret


# Several **metrics** can be used in computer vision depending on the objectives to reach

## ⚙️ Local Metrics

### ▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$



We set a threshold  $\varepsilon \in [0,1]$

Prediction are considered true if  $IoU \geq \varepsilon$ .

### ▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**


$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**

### ▪ Precision-recall curve :

Precision and recall are calculated for different confidence levels (proba of the class model) in order to obtain the curve:

$$\text{precision} = f(\text{recall})$$



## Model Metrics

### ▪ Mean Intersect over Union (mIoU)

For a fixed minimum confidence level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: over/under-prediction?
- Depends on a confidence threshold set



- Very "visual" metric, easy to interpret




# Several **metrics** can be used in computer vision depending on the objectives to reach

## ⚙️ Local Metrics

### ▪ Intersect over Union (IoU)

Measuring the common cover on an object can be formalized as:

$$IoU = \frac{\text{aire}(\text{Box}_{\text{True}} \cap \text{Box}_{\text{Predicted}})}{\text{aire}(\text{Box}_{\text{True}} \cup \text{Box}_{\text{Predicted}})}$$



We set a threshold  $\varepsilon \in [0,1]$

Prediction are considered true if  $IoU \geq \varepsilon$ .

### ▪ Precision/Recall :

$$\text{precision} = \frac{TP}{TP + FP}$$

Precision answers the question: **What share of the detected objects were the right ones?**


$$\text{recall} = \frac{TP}{TP + FN}$$

Recall answers the question: **What is the share of the objects that have been detected?**

### ▪ Precision-recall curve :

Precision and recall are calculated for different confidence levels (proba of the class model) in order to obtain the curve:

$$\text{precision} = f(\text{recall})$$



## Model Metrics

### ▪ Mean Intersect over Union (mIoU)

For a fixed **minimum confidence** level, the IoU is calculated for each annotation and then averaged over a class or globally.



- No indication of the type of error: on/under-prediction?
- Depends on a confidence threshold set

- Very "visual" metric, easy to interpret

### ▪ Average Precision (AP)

Once the IoU threshold is set, the AP of a class can be calculated as a measure of the area below the precision-recall curve.




- Difficult to interpret
- Depends on a fixed IoU threshold

- Independent of confidence level
- Gives an indication of the precision/recall trade-off

➤ **The mean Average Precision (mAP) is the AP averaged over all classes**

The AP reflects the ability to combine precision and recall for a model

### Example n°1




$$AP \approx 0.74$$

#### Interpretation :

- We see that we can maintain a good trade-off between the level of precision and the level of recall
  - The **green dot** guarantees an accuracy of about 80% for a 75% recall: most objects are spotted without too many errors
- An important AP ensures a certain efficiency of the model that will combine precision and recall

### Example n°2



$$AP \approx 0.28$$

#### Interpretation :

We see here that we cannot guarantee precision and recall: we must choose to favour one or the other

- The **green zone** guarantees a precision > 75% accuracy but also a recall < 20% : we don't make much mistake but we spot very few objects
  - The **pink zone** guarantees a recall > 80% but a precision < 5% : We spot almost all the object but most of the objects identified are not the one we want
- A low AP means it is difficult to reconcile accuracy and recall: one should be favored

**Resources:** all classical Machine Learning might come handy and a good understanding of Computer Vision libraries will be helpful

### Computer Vision libraries



*Standard library for  
deep learning*



*Standard library classical  
computer vision*



*An alternative to  
Pytorch*

### ML + Viz libraries



*To develop a wide  
range of ML models*



*To exploite model's  
output and  
aggregate them*



*To efficiently develop a  
dashboard / front-end*